

Gaussian processes and Gauss Markov random fields

Tudor Muntianu

December 3, 2020

Abstract

This report gives a (very) brief summary of work done this term while taking David Blei's course, Foundations of Graphical Models, as administered and supervised by Jeremy. The [GitHub repo](#) contains all of the code and notebooks that I'll reference here. Some of them are quite messy because they are still being tinkered with.

1 Motivation

The primary motivation for this independent study was the SuperEEG project, which is based largely on Gaussian process regression. Currently the model does not rely on much of the state-of-the-art for GPs largely due to the structure of the problem and the resulting difficulty in expressing it as a single covariance matrix. So the work from this term can be seen as an attempt to take a second crack at this problem while also gaining a (more) solid foundation in graphical models.

So rather than examine any dataset or problem (even ECoG data specifically) in particular, I wanted to explore and learn more about graphical models, GMRFs, GPs, VI, and the software packages that implement them, largely because a SuperEEG 2.0 would call for using something other than NumPy for computation.

2 The problem

The basic idea of SuperEEG is to impute spatial data at locally unobserved but globally observed locations, and using this global structure to improve these imputed values. There are not enough papers on this subject, and many of them come from the field of geostatistics where locations are never observed globally. In ECoG at some point every location is sampled, but that is not the case on Earth. This is why GPs can be used in geostatistics, see Alvarez's review [1] which contains some examples, and why SuperEEG's flavor of GPs is novel.

There are, however, some developments which have bled over from computer vision. Deep GMRFs [8] and other inpainting models (which are really just imputing values in images like deep image priors [10] offer some interesting perspectives on this sort of spatial imputation. So what I did here was 1) implement a relatively simple but very powerful GP architecture to learn some variational inference and some other GP-specific methods such as inducing point methods, and 2) reimplement the Siden and Lindsten's deep GMRF model because their code is... neither particularly extensible nor readable. I also picked this model because there are a number of clear future directions that are not as difficult as in the world of theoretical GPs (which is astonishingly deep and convoluted).

3 Image classification and GPs

For the GP part of the term, I chose to implement and play with van der Wilk’s convolutional Gaussian process. van der Wilk’s paper is significant because of the historical Gaussian process struggle with images. This struggle originates from two facts: 1) GPs are $\mathcal{O}(n^3)$ for training, and require $\mathcal{O}(n^2)$ memory, making dealing with thousands of images times thousands of pixels difficult, and 2) many traditional/commonly used kernels are stationary (translation invariant), making it difficult for them to capture image structure in the way CNNs and other deep architectures can. These limitations are not unique to images, but with MNIST and CIFAR being common datasets, the limitations of GPs are quickly seen when applied to these. However, variational inference and sparse approximations have done much to alleviate the first issue, and newer state-of-the-art methods like deep kernel learning [12] and deep convolutional GPs [3] have been tackling the second.

The way shallow convolutional GPs attack this problem is in much the same way. The model uses inducing points and stochastic VI to make training and inference possible, and the structure of the convolutional kernel allows the stationary radial basis function to also acquire certain non-stationary properties by operating on individual patches and taking weighted sums over them. The exact mechanism of this operation will be described later, but intuitively, it allows the kernel to express relationships between distant pixels while also encoding information about proximal pixels through the RBF. It’s a useful (albeit at this point dated) extension of CNNs to GPs much in the same way Siden’s deep GMRF paper does the same thing for GMRFs.

4 Deep GMRFs and inpainting

A brief introduction to GMRFs is simple: a GMRF is just a random vector from a multivariate Gaussian that satisfies certain conditional independence assumptions. These assumptions can be modified depending on the context, but are always present (hence Gauss Markov). See Rue and Held for more [7]. These models have been largely ignored as they relied largely on the advantage that the representing precision matrices are sparse, and so clever numerical algorithms to operate efficiently on these matrices yielded good results in the past. Obviously this advantage diminishes daily now that we have great autodiff and better/easier sampling tools.

Nevertheless GMRFs are probabilistic graphical models that can carry forward uncertainty estimates, while also benefiting from VI techniques, so they still have utility albeit not in their most rudimentary forms.

The deep GMRF model is simple, though interestingly it defines the generative process “backwards”. The process is as follows (ignoring many details of the CNN), given an image with missing pixels:

1. Sample from a GMRF with the image as a mean, and $\sigma^2 \mathbf{I}$ as the covariance, where σ^2 is a learned parameter
2. Put this sample through a (special) CNN such that the output is a GMRF distributed according to $\mathcal{N}(\mathbf{0}, \mathbf{I})$

Intuitively then this process is just sampling a standard normal GMRF and then sending it backwards through a CNN to get the image, though in this case the latent variables we are learning are the pixels of the image regardless of whether they are missing.

5 Quick summary of current state of notebooks

The GP notebooks are readable and pretty self-explanatory, so I’m not going to discuss them much here. I also think the GMRF direction is more interesting anyway.

There is a `convgps` folder in the GP notebook folder that contains the GPyTorch implementation of convolutional kernels.

The notebooks are a bit of a mess since I’m still breaking everything. Here are the results so far.




Figure 1: From left to right: original image, masked image, posterior VI mean

The results of the VI mean are surprisingly good, especially given what the Siden paper reports. They use a diagonal multivariate Gaussian as the variational distribution, and admit that this is a very restrictive distribution. Their reasoning is that they use VI only to learn the model parameters, and then they use conjugate gradient to actually sample the posterior. This works because the CNN constraints they describe requires that the transformation be full rank in order to produce a valid GMRF. Then you can just pretend that the entire CNN is one matrix and perform CG to get a solution, which gives a way of getting the posterior mean and variance, as well as allowing sampling.

I’ve included my own CG implementation in the notebook but have not adapted it to the CNN mostly because this method is extremely inelegant. Using VI just to optimize model parameters and then using a linear system solver (on what is ultimately a non-linear transform if using ReLU or any non-linear activation function) to sample the posterior is redundant. The authors admit this, but still rather than replicate the same hack they did I’d rather try to improve the variational distribution.

Regardless, the image above is just a bunch of sine and cosine functions (see notebooks for details), and part of the image is masked. Using the same diagonal multivariate Gaussian variational distribution, though, the model still managed to achieve impressive inpainting, although it added significant noise to the image. Rather than play around with more parameters to reduce this noise, though, I’m currently trying to come up with a more expressive variational distribution, especially since Pyro allows very flexible creations of variational distribution or “guides.”

Another really interesting observation is that the missing pixels are learned earlier than the observed pixels.




Figure 2: Variational posterior mean at iteration 1000 of training

This indicates again that the variational distribution is the limiting factor, since it constrains the existing pixels rather than the missing ones, while the model itself is able to quickly learn the structure and begin imputing well even after limited training.

Apart from coming up with a better variational distribution, there's some more changes that might be made. The current model only operates on one image, though since the CNN parameters hopefully learn the global structure, we could treat each image as conditionally independent and then optimize these parameters (or a superset of them created to better deal with multiple images) in order to train on different sets of imputed pixels.

So these are the two areas for more work: 1) better variational distribution, 2) extend the model to allow conditional independence of images with different masks or missing pixels.

6 Quick note on GPyTorch and GPFlow

When implementing the conv GPs, I looked at van der Wilk's implementations in GPFlow, a TensorFlow based GP library. I intentionally chose GPyTorch to make it harder for myself (no cheating or copy-paste), but also so that I could compare the two.

While I prefer GPFlow due to the sheer readability of the existing code, and the large array of models already implemented, I think GPyTorch has a more powerful stack. Integration with Pyro makes MCMC and VI stupid easy (and Stan borderline irrelevant), more sophisticated lazy and approximate covariance calculations are powerful but difficult to leverage for custom models, although the developers respond to requests for help or features quickly. GPyTorch is also seeing significant development activity and funding, so I would expect it to continue to outpace GPFlow from a technical perspective, but perhaps not from a usability perspective.

Both are crucially missing basic explanations of how and where inference lives in the model, kernel, and likelihood objects. This makes building truly custom models from scratch extremely painful the first time around. I spent at least five hours just reading through the codebases in an

attempt to decipher the package’s architecture. Pyro is much the same way. Once you need to implement a custom loss function or change the way the ELBO is computed docs get thin and examples are not found.

References

- [1] Mauricio A. Alvarez, Lorenzo Rosasco, and Neil D. Lawrence. Kernels for vector-valued functions: a review, 2012.
- [2] David M. Blei, Alp Kucukelbir, and Jon D. McAuliffe. Variational inference: A review for statisticians. *Journal of the American Statistical Association*, 112(518):859–877, Apr 2017.
- [3] Vincent Dutordoir, Mark van der Wilk, Artem Artemev, and James Hensman. Bayesian image classification with deep convolutional gaussian processes, 2019.
- [4] James Hensman, Nicolo Fusi, and Neil D. Lawrence. Gaussian processes for big data, 2013.
- [5] Lucy L W Owen, Tudor A Muntianu, Andrew C Heusser, Patrick M Daly, Katherine W Scangos, and Jeremy R Manning. A Gaussian Process Model of Human Electrocorticographic Data. *Cerebral Cortex*, 06 2020. bhaa115.
- [6] CE. Rasmussen and CKI. Williams. *Gaussian Processes for Machine Learning*. Adaptive Computation and Machine Learning. MIT Press, Cambridge, MA, USA, January 2006.
- [7] Havard Rue and Leonhard Held. *Gaussian Markov Random Fields: Theory and Applications*. Monographs on Statistics and Applied Probability. Chapman & Hall/CRC, 2005.
- [8] Per Sidén and Fredrik Lindsten. Deep gaussian markov random fields, 2020.
- [9] Michalis Titsias. Variational learning of inducing variables in sparse gaussian processes. volume 5 of *Proceedings of Machine Learning Research*, pages 567–574, Hilton Clearwater Beach Resort, Clearwater Beach, Florida USA, 16–18 Apr 2009. PMLR.
- [10] Dmitry Ulyanov, Andrea Vedaldi, and Victor Lempitsky. Deep image prior. *International Journal of Computer Vision*, 128(7):1867–1888, Mar 2020.
- [11] Mark van der Wilk, Carl Edward Rasmussen, and James Hensman. Convolutional gaussian processes, 2017.
- [12] Andrew Gordon Wilson, Zhiting Hu, Ruslan Salakhutdinov, and Eric P. Xing. Deep kernel learning, 2015.